

finalProject

June 11, 2020

1 EDA explosive data analysis

```
[0]: import pandas as pd
labels = pd.read_csv('labels.csv')
```

```
[0]: n = len(labels)
breed = set(labels['breed'])
n_class = len(breed)
class_to_num = dict(zip(breed, range(n_class)))
n_class = len(breed)
```

```
[0]: labels.head(5)
```

```
[0]:
```

	id	breed
0	000bec180eb18c7604dcecc8fe0dba07	boston_bull
1	001513dfcb2ffafc82cccf4d8bbaba97	dingo
2	001cdf01b096e06d78e9e5112d419397	pekinese
3	00214f311d5d2247d5dfe4fe24b2303d	bluetick
4	0021f9ceb3235effd7fcde7f7538ed62	golden_retriever

```
[0]: yy = pd.value_counts(labels['breed'])
print(yy[:10])
```

```
scottish_deerhound    126
maltese_dog           117
afghan_hound          116
entlebucher           115
bernese_mountain_dog  114
shih-tzu              112
great_pyrenees        111
pomeranian            111
basenji               110
samoyed               109
Name: breed, dtype: int64
```

```
[0]: print(yy[-10:])
```

```


otterhound          69
giant_schnauzer     69
tibetan_mastiff     69
german_shepherd     69
walker_hound        69
golden_retriever    67
komondor            67
brabancon_griffon   67
eskimo_dog          66
briard              66
Name: breed, dtype: int64

```

```
[0]: import matplotlib.pyplot as plt
import seaborn as sns
fig, ax = plt.subplots()
fig.set_size_inches(20,10)
sns.set_style("whitegrid")

ax = sns.barplot(x = yy.index, y = yy, data = labels)
ax.set_xticklabels(ax.get_xticklabels(), rotation = 90, fontsize = 10)
ax.set(xlabel='Dog Breed', ylabel='Count')
ax.set_title('Distribution of the Dog Breeds')
```

```
[0]: Text(0.5, 1.0, 'Distribution of the Dog Breeds')
```



2 Resize the data to be three new data set as 120 * 20, 120 * 40 and 120 * 60

```
[0]: !mkdir images120_20
      !mkdir images120_40
      !mkdir images120_60
```

```
[0]: import os, random, shutil

for dir in os.listdir('Images/'):
    os.mkdir('images120_20/'+dir)
    os.mkdir('images120_40/'+dir)
    os.mkdir('images120_60/'+dir)
```

```
[0]: # randomly select 20/ 40 / 60 from the original dataset of 120 do breeds

for dir in os.listdir('Images/'):
    filepath = 'Images/' + dir + '/'
    tarDir = 'images120_20/' + dir + '/'

    files = os.listdir(filepath)
    sample = random.sample(files, 20)

    for name in sample:
        shutil.copy(filepath+name, tarDir+name)
```

```
[0]: # randomly select 20/ 40 / 60 from the original dataset of 120 do breeds

for dir in os.listdir('Images/'):
    filepath = 'Images/' + dir + '/'
    tarDir = 'images120_40/' + dir + '/'

    files = os.listdir(filepath)
    sample = random.sample(files, 40)

    for name in sample:
        shutil.copy(filepath+name, tarDir+name)
```

```
[0]: # randomly select 20/ 40 / 60 from the original dataset of 120 do breeds

for dir in os.listdir('Images/'):
    filepath = 'Images/' + dir + '/'
    tarDir = 'images120_60/' + dir + '/'

    files = os.listdir(filepath)
    sample = random.sample(files, 60)
```

```
for name in sample:
    shutil.copy(filepath+name, tarDir+name)
```

```
[0]: !mkdir images120_60_train
!mkdir images120_60_val
!mkdir images120_60_test
```

```
[0]: for dir in os.listdir('Images/'):
    os.mkdir('images120_60_train/'+dir)
    os.mkdir('images120_60_test/'+dir)
    os.mkdir('images120_60_val/'+dir)
```

```
[0]: for dir in os.listdir('Images/'):
    filepath = 'Images/' + dir + '/'
    files = os.listdir(filepath)
    sample = random.sample(files, 60)

    tarDir = 'images120_60_train/' + dir + '/'
    for name in sample[:40]:
        shutil.copy(filepath+name, tarDir+name)

    tarDir = 'images120_60_val/' + dir + '/'
    for name in sample[40:50]:
        shutil.copy(filepath+name, tarDir+name)

    tarDir = 'images120_60_test/' + dir + '/'
    for name in sample[50:]:
        shutil.copy(filepath+name, tarDir+name)
```

3 training process starts from here

4 Pytorch

```
[5]: from __future__ import print_function, division

import torch
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
import torchvision
from torchvision import datasets, models, transforms
```

```
import matplotlib.pyplot as plt
import numpy as np
import time
import os
import copy
```

```
[6]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
```

```
[6]: device(type='cuda')
```

```
[7]: mean_nums = [0.485, 0.456, 0.406]
std_nums = [0.229, 0.224, 0.225]

chosen_transforms = {'train': transforms.Compose([
    transforms.RandomResizedCrop(size=256),
    transforms.RandomRotation(degrees=15),
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(mean_nums, std_nums)
]), 'val': transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean_nums, std_nums)
]), 'test': transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean_nums, std_nums)
])
}
```

```
[8]: # Set the directory for the data
data_dir = 'images120_60_'

# Use the image folder function to create datasets
dog_datasets = {x: datasets.ImageFolder(data_dir+x+'/',
    chosen_transforms[x]) for x in ['train', 'val', 'test']}
```

```
[9]: dataloaders = {x: DataLoader(dataset=dog_datasets[x], batch_size=8,
    ↪shuffle=True, num_workers=4)
    for x in ['train', 'val', 'test']}
```

```
[10]: dataset_sizes = {x: len(dog_datasets[x]) for x in ['train', 'val', 'test']}
class_names = dog_datasets['train'].classes
```

```
print(dataset_sizes)
print(len(class_names))
```

```
{'train': 4800, 'val': 1200, 'test': 1200}
120
```

```
[11]: # Visualize some images
def imshow(inp, title=None):
    inp = inp.numpy().transpose((1, 2, 0))
    mean = np.array([mean_nums])
    std = np.array([std_nums])
    inp = std * inp + mean
    inp = np.clip(inp, 0, 1)
    plt.imshow(inp)
    if title is not None:
        plt.title(title)
    plt.pause(0.001) # Pause a bit so that plots are updated

# Grab some of the training data to visualize
inputs, classes = next(iter(dataloaders['train']))

# Now we construct a grid from batch
out = torchvision.utils.make_grid(inputs)

imshow(out, title=[class_names[x] for x in classes])
```

['n02107683-Bernese_mountain_dog', 'n02116738-African_hunting_dog', 'n02104029-kuyasr', 'n02095889-Sealyham_terrier', 'n02102973-Irish_water_spaniel', 'n02090379-redbone', 'n02113023-Pembroke', 'n02095570-Lakeland_terrier']



```
[12]: LOG_DIR = 'log/'
def train_model(model, criterion, optimizer, num_epochs=10,
    log_file_name='log'):
    log_file = open(LOG_DIR + log_file_name, 'w')
    since = time.time()

    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    for epoch in range(num_epochs):
        print('Epoch {}/{}'.format(epoch, num_epochs - 1))
        print('-' * 10)

        # Each epoch has a training and validation phase
        for phase in ['train', 'val']:
            if phase == 'train':
```

```

        # scheduler.step()
        model.train() # Set model to training mode
    else:
        model.eval() # Set model to evaluate mode

    current_loss = 0.0
    current_corrects = 0

    # Here's where the training happens
    # print('Iterating through data...')

    for inputs, labels in dataloaders[phase]:
        inputs = inputs.to(device)
        labels = labels.to(device)

        # We need to zero the gradients, don't forget it
        optimizer.zero_grad()

        outputs = model(inputs)
        _, preds = torch.max(outputs, 1)
        loss = criterion(outputs, labels)

        # backward + optimize only if in training phase
        if phase == 'train':
            loss.backward()
            optimizer.step()

        # We want variables to hold the loss statistics
        current_loss += loss.item() # * inputs.size(0)
        current_corrects += torch.sum(preds == labels.data)

    epoch_loss = current_loss / dataset_sizes[phase]
    epoch_acc = current_corrects.double() / dataset_sizes[phase]

    outline = f'{epoch} {epoch_loss:.4f} {epoch_acc:.4f} {phase}\n'
    log_file.write(outline)

    print('{} Loss: {:.4f} Acc: {:.4f}'.format(
        phase, epoch_loss, epoch_acc))

    # Make a copy of the model if the accuracy on the validation set
    ↪ has improved
    if phase == 'val' and epoch_acc > best_acc:
        best_acc = epoch_acc
        best_model_wts = copy.deepcopy(model.state_dict())

    print()

```

```

log_file.close()
time_since = time.time() - since
print('Training complete in {:.0f}m {:.0f}s'.format(
    time_since // 60, time_since % 60))
print('Best val Acc: {:.4f}'.format(best_acc))

# Now we'll load in the best model weights and return it
model.load_state_dict(best_model_wts)
return model

```

```

[13]: def evaluate_model(model, log_file_name='log'):

    phase = 'test'
    log_file = open(LOG_DIR + log_file_name, 'a')
    model.eval() # Set model to evaluate mode

    current_loss = 0.0
    current_corrects = 0

    for inputs, labels in dataloaders[phase]:
        inputs = inputs.to(device)
        labels = labels.to(device)

        outputs = model(inputs)
        _, preds = torch.max(outputs, 1)
        loss = criterion(outputs, labels)

        # We want variables to hold the loss statistics
        current_loss += loss.item() # * inputs.size(0)
        current_corrects += torch.sum(preds == labels.data)

    epoch_loss = current_loss / dataset_sizes[phase]
    epoch_acc = current_corrects.double() / dataset_sizes[phase]

    outline = f'{phase} {epoch_loss:.4f} {epoch_acc:.4f} {phase}\n'
    log_file.write(outline)
    log_file.close()

    print('{} Loss: {:.4f} Acc: {:.4f}'.format(
        phase, epoch_loss, epoch_acc))

```


5 Model 1– CNN

```
[14]: import torch.nn.functional as F

n_classes = 120

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()

        ## Define layers of a CNN
        # convolutional layer (sees 224x224x3 image tensor)
        self.conv1 = nn.Conv2d(3, 16, 3, padding = 1 )
        # convolutional layer (sees 112x112x16 image tensor)
        self.conv2 = nn.Conv2d(16, 32, 3, padding = 1)
        # convolutional layer (sees 56x56x32 image tensor)
        self.conv3 = nn.Conv2d(32, 64, 3, padding = 1)
        # convolutional layer (sees 28x28x64 image tensor)
        self.conv4 = nn.Conv2d(64, 128, 3, padding = 1)
        # convolutional layer (sees 14x14x128 image tensor)
        self.conv5 = nn.Conv2d(128, 256, 3, padding = 1 )

        # max pooling layer
        self.pool = nn.MaxPool2d(2, 2)
        # dropout layer (p=0.2)
        self.dropout = nn.Dropout(0.2)

        self.conv_bn1 = nn.BatchNorm2d(224,3)
        self.conv_bn2 = nn.BatchNorm2d(16)
        self.conv_bn3 = nn.BatchNorm2d(32)
        self.conv_bn4 = nn.BatchNorm2d(64)
        self.conv_bn5 = nn.BatchNorm2d(128)
        self.conv_bn6 = nn.BatchNorm2d(256)

        # linear layer (256 * 7 * 7 -> 512)
        self.fc1 = nn.Linear(256 * 7 * 7, 512)
        # linear layer (256 * 7 * 7 -> n_classes (133))
        self.fc2 = nn.Linear(512, n_classes)

    def forward(self, x):
        ## Define forward behavior
        # add sequence of convolutional and max pooling layers
        x = self.pool(F.relu(self.conv1(x)))
        x = self.conv_bn2(x)
        x = self.pool(F.relu(self.conv2(x)))
```

```

x = self.conv_bn3(x)
x = self.pool(F.relu(self.conv3(x)))
x = self.conv_bn4(x)
x = self.pool(F.relu(self.conv4(x)))
x = self.conv_bn5(x)
x = self.pool(F.relu(self.conv5(x)))
x = self.conv_bn6(x)

# flatten image input
x = x.view(-1, 256 * 7 * 7)
# add dropout layer
x = self.dropout(x)
# add second hidden layer
x = F.relu(self.fc1(x))
x = self.dropout(x)
x = self.fc2(x)
return x

```

```

[ ]: cnn = Net()
cnn = cnn.to(device)
criterion = nn.CrossEntropyLoss().to(device)
optimizer_ft = optim.SGD(cnn.parameters(), lr=0.00001, momentum=0.9)

cnn_best = train_model(cnn, criterion, optimizer_ft, num_epochs=10,
                       log_file_name='cnn_60_0dot00001')

print('-' * 10)
evaluate_model(cnn_best, log_file_name='cnn_60_0dot00001')

```

Epoch 0/9

train Loss: 0.6043 Acc: 0.0075
val Loss: 0.6044 Acc: 0.0067

Epoch 1/9

train Loss: 0.6025 Acc: 0.0121
val Loss: 0.6017 Acc: 0.0067

Epoch 2/9

train Loss: 0.6007 Acc: 0.0085
val Loss: 0.5995 Acc: 0.0075

Epoch 3/9

train Loss: 0.5989 Acc: 0.0138

val Loss: 0.5979 Acc: 0.0117

Epoch 4/9

train Loss: 0.5985 Acc: 0.0158

val Loss: 0.5959 Acc: 0.0100

Epoch 5/9

train Loss: 0.5965 Acc: 0.0158

6 Model2: Resnet34

```
[0]: resnet = models.resnet34(pretrained=True)
      # for param in resnet.parameters():
      #     param.requires_grad = False

      num_fts = resnet.fc.in_features
      resnet.fc = nn.Linear(num_fts, 120)

      # lr = [0.001, 0.0001, 0.00001]
      # log_file_name = ['resnet_40_0dot001', 'resnet_40_0dot0001',
      #                  ↪ 'resnet_40_0dot00001']

      resnet = resnet.to(device)
      criterion = nn.CrossEntropyLoss().to(device)
      optimizer_ft = optim.SGD(resnet.parameters(), lr=0.0001, momentum=0.9)

      resnet_best = train_model(resnet, criterion, optimizer_ft,
                                num_epochs=20, log_file_name='resnet_60_0dot0001')

      print('-' * 10)
      evaluate_model(resnet_best, log_file_name='resnet_60_0dot0001')
```

Downloading: "https://download.pytorch.org/models/resnet34-333f7ec4.pth" to
/root/.cache/torch/checkpoints/resnet34-333f7ec4.pth

HBox(children=(FloatProgress(value=0.0, max=87306240.0), HTML(value='')))

Epoch 0/19

train Loss: 0.6009 Acc: 0.0204

val Loss: 0.5626 Acc: 0.0483

Epoch 1/19

train Loss: 0.5658 Acc: 0.0577
val Loss: 0.4914 Acc: 0.1733

Epoch 2/19

train Loss: 0.5294 Acc: 0.1250
val Loss: 0.4136 Acc: 0.3217

Epoch 3/19

train Loss: 0.4920 Acc: 0.1854
val Loss: 0.3448 Acc: 0.4467

Epoch 4/19

train Loss: 0.4589 Acc: 0.2331
val Loss: 0.2832 Acc: 0.5400

Epoch 5/19

train Loss: 0.4325 Acc: 0.2827
val Loss: 0.2421 Acc: 0.5883

Epoch 6/19

train Loss: 0.4102 Acc: 0.2981
val Loss: 0.2087 Acc: 0.6267

Epoch 7/19

train Loss: 0.3923 Acc: 0.3346
val Loss: 0.1833 Acc: 0.6592

Epoch 8/19

train Loss: 0.3733 Acc: 0.3533
val Loss: 0.1707 Acc: 0.6775

Epoch 9/19

train Loss: 0.3593 Acc: 0.3744
val Loss: 0.1581 Acc: 0.6917

Epoch 10/19

train Loss: 0.3501 Acc: 0.3858
val Loss: 0.1418 Acc: 0.7150

Epoch 11/19

train Loss: 0.3396 Acc: 0.3996

val Loss: 0.1371 Acc: 0.7258

Epoch 12/19

train Loss: 0.3318 Acc: 0.4123

val Loss: 0.1309 Acc: 0.7292

Epoch 13/19

train Loss: 0.3233 Acc: 0.4173

val Loss: 0.1215 Acc: 0.7250

Epoch 14/19

train Loss: 0.3138 Acc: 0.4383

val Loss: 0.1197 Acc: 0.7333

Epoch 15/19

train Loss: 0.3106 Acc: 0.4294

val Loss: 0.1189 Acc: 0.7333

Epoch 16/19

train Loss: 0.3037 Acc: 0.4438

val Loss: 0.1097 Acc: 0.7583

Epoch 17/19

train Loss: 0.2946 Acc: 0.4590

val Loss: 0.1098 Acc: 0.7475

Epoch 18/19

train Loss: 0.2959 Acc: 0.4471

val Loss: 0.1050 Acc: 0.7642

Epoch 19/19

train Loss: 0.2889 Acc: 0.4646

val Loss: 0.1016 Acc: 0.7592

Training complete in 15m 31s

Best val Acc: 0.764167

test Loss: 0.1053 Acc: 0.7333

```
[0]: def cm(model):
      nb_classes = 120

      confusion_matrix = torch.zeros(nb_classes, nb_classes)
      with torch.no_grad():
          for i, (inputs, classes) in enumerate(dataloaders['test']):
              inputs = inputs.to(device)
              classes = classes.to(device)
              outputs = model(inputs)
              _, preds = torch.max(outputs, 1)
              for t, p in zip(classes.view(-1), preds.view(-1)):
                  confusion_matrix[t.long(), p.long()] += 1

      return confusion_matrix

cm_resnet = cm(resnet_best)
cm_resnet
```

```
[0]: tensor([[ 9.,  0.,  0., ...,  0.,  0.,  0.],
            [ 0.,  7.,  0., ...,  0.,  0.,  0.],
            [ 0.,  0., 10., ...,  0.,  0.,  0.],
            ...,
            [ 0.,  0.,  0., ..., 10.,  0.,  0.],
            [ 0.,  0.,  0., ...,  0., 10.,  0.],
            [ 0.,  0.,  0., ...,  0.,  0., 10.]])
```

```
[0]: cm_resnet.shape
```

```
[0]: torch.Size([120, 120])
```

```
[0]: my_class_names = [x.split('-')[1] for x in class_names]
      my_class_names
```


```
[0]: ['Chihuahua',
      'Japanese_spaniel',
      'Maltese_dog',
      'Pekinese',
      'Shih',
      'Blenheim_spaniel',
      'papillon',
      'toy_terrier',
      'Rhodesian_ridgeback',
      'Afghan_hound',
      'basset',
```

'beagle',
'bloodhound',
'bluetick',
'black',
'Walker_hound',
'English_foxhound',
'redbone',
'borzoi',
'Irish_wolfhound',
'Italian_greyhound',
'whippet',
'Ibizan_hound',
'Norwegian_elkhound',
'otterhound',
'Saluki',
'Scottish_deerhound',
'Weimaraner',
'Staffordshire_bullterrier',
'American_Staffordshire_terrier',
'Bedlington_terrier',
'Border_terrier',
'Kerry_blue_terrier',
'Irish_terrier',
'Norfolk_terrier',
'Norwich_terrier',
'Yorkshire_terrier',
'wire',
'Lakeland_terrier',
'Sealyham_terrier',
'Airedale',
'cairn',
'Australian_terrier',
'Dandie_Dinmont',
'Boston_bull',
'miniature_schnauzer',
'giant_schnauzer',
'standard_schnauzer',
'Scotch_terrier',
'Tibetan_terrier',
'silky_terrier',
'soft',
'West_Highland_white_terrier',
'Lhasa',
'flat',
'curly',
'golden_retriever',
'Labrador_retriever',

'Chesapeake_Bay_retriever',
'German_short',
'vizsla',
'English_setter',
'Irish_setter',
'Gordon_setter',
'Brittany_spaniel',
'clumber',
'English_springer',
'Welsh_springer_spaniel',
'cocker_spaniel',
'Sussex_spaniel',
'Irish_water_spaniel',
'kuvasz',
'schipperke',
'groenendael',
'malinois',
'briard',
'kelpie',
'komondor',
'Old_English_sheepdog',
'Shetland_sheepdog',
'collie',
'Border_collie',
'Bouvier_des_Flandres',
'Rottweiler',
'German_shepherd',
'Doberman',
'miniature_pinscher',
'Greater_Swiss_Mountain_dog',
'Bernese_mountain_dog',
'Appenzeller',
'EntleBucher',
'boxer',
'bull_mastiff',
'Tibetan_mastiff',
'French_bulldog',
'Great_Dane',
'Saint_Bernard',
'Eskimo_dog',
'malamute',
'Siberian_husky',
'affenpinscher',
'basenji',
'pug',
'Leonberg',
'Newfoundland',


```
'Great_Pyrenees',  
'Samoyed',  
'Pomeranian',  
'chow',  
'keeshond',  
'Brabancon_griffon',  
'Pembroke',  
'Cardigan',  
'toy_poodle',  
'miniature_poodle',  
'standard_poodle',  
'Mexican_hairless',  
'dingo',  
'dhole',  
'African_hunting_dog']
```

```
[0]: import seaborn as sn  
import pandas as pd  
# cm_resnet = cm(resnet_best)  
df_cm = pd.DataFrame(cm_resnet.numpy(), index = [i for i in my_class_names],  
                     columns = [i for i in my_class_names])  
plt.figure(figsize = (30,30))  
sn.heatmap(df_cm)  
plt.savefig('confusion_matrix')
```



7 Model3: VGG19

```
[0]: vgg = models.vgg19(pretrained=True)

vgg.classifier[6] = nn.Linear(4096, 120)

# lr = [0.001, 0.0001, 0.00001]
# log_file_name = ['vgg_60_0dot001', 'vgg_60_0dot0001', 'vgg_60_0dot00001']
```

```

vgg = vgg.to(device)
criterion = nn.CrossEntropyLoss().to(device)
optimizer_ft = optim.SGD(vgg.parameters(), lr=0.0001, momentum=0.9)

vgg_best = train_model(vgg, criterion, optimizer_ft,
                        num_epochs=20, log_file_name='vgg_20_0dot0001')

print('-' * 10)
evaluate_model(vgg_best, log_file_name='vgg_20_0dot0001')

```

Epoch 0/19

train Loss: 0.5927 Acc: 0.0167

val Loss: 0.5408 Acc: 0.1333

Epoch 1/19

train Loss: 0.5320 Acc: 0.0988

val Loss: 0.2870 Acc: 0.4417

Epoch 2/19

train Loss: 0.4148 Acc: 0.2393

val Loss: 0.1736 Acc: 0.5778

Epoch 3/19

train Loss: 0.3568 Acc: 0.3054

val Loss: 0.1428 Acc: 0.6500

Epoch 4/19

train Loss: 0.3209 Acc: 0.3744

val Loss: 0.1286 Acc: 0.6833

Epoch 5/19

train Loss: 0.3009 Acc: 0.3917

val Loss: 0.1225 Acc: 0.6917

Epoch 6/19

train Loss: 0.2930 Acc: 0.4089

val Loss: 0.1141 Acc: 0.7056

Epoch 7/19

train Loss: 0.2829 Acc: 0.4411
val Loss: 0.1129 Acc: 0.7111

Epoch 8/19

train Loss: 0.2734 Acc: 0.4351
val Loss: 0.1039 Acc: 0.7389

Epoch 9/19

train Loss: 0.2694 Acc: 0.4411
val Loss: 0.1073 Acc: 0.7222

Epoch 10/19

train Loss: 0.2533 Acc: 0.4750
val Loss: 0.1084 Acc: 0.7278

Epoch 11/19

train Loss: 0.2462 Acc: 0.4899
val Loss: 0.1098 Acc: 0.7306

Epoch 12/19

train Loss: 0.2393 Acc: 0.5024
val Loss: 0.0995 Acc: 0.7472

Epoch 13/19

train Loss: 0.2378 Acc: 0.5131
val Loss: 0.1129 Acc: 0.7139

Epoch 14/19

train Loss: 0.2447 Acc: 0.4893
val Loss: 0.1138 Acc: 0.7222

Epoch 15/19

train Loss: 0.2320 Acc: 0.5196
val Loss: 0.1019 Acc: 0.7417

Epoch 16/19

train Loss: 0.2264 Acc: 0.5280
val Loss: 0.1097 Acc: 0.7333

```

Epoch 17/19
-----
train Loss: 0.2270 Acc: 0.5304
val Loss: 0.1069 Acc: 0.7389

Epoch 18/19
-----
train Loss: 0.2234 Acc: 0.5167
val Loss: 0.1131 Acc: 0.7139

Epoch 19/19
-----
train Loss: 0.2265 Acc: 0.5238
val Loss: 0.1054 Acc: 0.7500

Training complete in 16m 6s
Best val Acc: 0.750000
-----
test Loss: 0.1005 Acc: 0.7861

```

8 Model4: Inception V3

```

[0]: mean_nums = [0.485, 0.456, 0.406]
     std_nums = [0.229, 0.224, 0.225]

     chosen_transforms = {'train': transforms.Compose([
         transforms.RandomResizedCrop(size=324),
         transforms.RandomRotation(degrees=15),
         transforms.RandomResizedCrop(299),
         transforms.RandomHorizontalFlip(),
         transforms.ToTensor(),
         transforms.Normalize(mean_nums, std_nums)
     ]), 'val': transforms.Compose([
         transforms.Resize(324),
         transforms.CenterCrop(299),
         transforms.ToTensor(),
         transforms.Normalize(mean_nums, std_nums)
     ]), 'test': transforms.Compose([
         transforms.Resize(324),
         transforms.CenterCrop(299),
         transforms.ToTensor(),
         transforms.Normalize(mean_nums, std_nums)
     ])
     }

     # Set the directory for the data

```

```

data_dir = 'images120_60_'

# Use the image folder function to create datasets
dog_datasets = {x: datasets.ImageFolder(data_dir+x+'/',
    chosen_transforms[x]) for x in ['train', 'val', 'test']}

dataloaders = {x: DataLoader(dataset=dog_datasets[x], batch_size=8,
    ↪shuffle=True, num_workers=4)
    for x in ['train', 'val', 'test']}

dataset_sizes = {x: len(dog_datasets[x]) for x in ['train', 'val', 'test']}
class_names = dog_datasets['train'].classes
print(dataset_sizes)
print(len(class_names))

```

```

{'train': 4800, 'val': 1200, 'test': 1200}
120

```

```

[0]: LOG_DIR = 'log/'
def train_model_inception(model, criterion, optimizer, num_epochs=10,
    ↪log_file_name='log'):
    log_file = open(LOG_DIR + log_file_name, 'w')
    since = time.time()

    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    for epoch in range(num_epochs):
        print('Epoch {}/{}'.format(epoch, num_epochs - 1))
        print('-' * 10)

        # Each epoch has a training and validation phase
        for phase in ['train', 'val']:
            if phase == 'train':
                # scheduler.step()
                model.train() # Set model to training mode
            else:
                model.eval() # Set model to evaluate mode

            current_loss = 0.0
            current_corrects = 0

            # Here's where the training happens
            # print('Iterating through data...')

            for inputs, labels in dataloaders[phase]:
                inputs = inputs.to(device)

```

```

labels = labels.to(device)

# We need to zero the gradients, don't forget it
optimizer.zero_grad()

if phase == 'train':
    outputs, aux_outputs = model(inputs)
    loss1 = criterion(outputs, labels)
    loss2 = criterion(aux_outputs, labels)
    loss = loss1 + 0.4 * loss2
else:
    outputs = model(inputs)
    loss = criterion(outputs, labels)

_, preds = torch.max(outputs, 1)

# backward + optimize only if in training phase
if phase == 'train':
    loss.backward()
    optimizer.step()

# We want variables to hold the loss statistics
current_loss += loss.item() # * inputs.size(0)
current_corrects += torch.sum(preds == labels.data)

epoch_loss = current_loss / dataset_sizes[phase]
epoch_acc = current_corrects.double() / dataset_sizes[phase]

outline = f'{epoch} {epoch_loss:.4f} {epoch_acc:.4f} {phase}\n'
log_file.write(outline)

print('{} Loss: {:.4f} Acc: {:.4f}'.format(
    phase, epoch_loss, epoch_acc))

# Make a copy of the model if the accuracy on the validation set
↳ has improved
if phase == 'val' and epoch_acc > best_acc:
    best_acc = epoch_acc
    best_model_wts = copy.deepcopy(model.state_dict())

print()

log_file.close()
time_since = time.time() - since
print('Training complete in {:.0f}m {:.0f}s'.format(
    time_since // 60, time_since % 60))

```

```

print('Best val Acc: {:.4f}'.format(best_acc))

# Now we'll load in the best model weights and return it
model.load_state_dict(best_model_wts)
return model

```

```

[0]: inception = models.inception_v3(pretrained=True)

inception.AuxLogits.fc = nn.Linear(768, 120)
inception.fc = nn.Linear(2048, 120)

# lr = [0.001, 0.0001, 0.00001]
# log_file_name = ['inception_60_0dot001', 'inception_60_0dot0001',
    ↳ 'inception_60_0dot00001']

inception = inception.to(device)
criterion = nn.CrossEntropyLoss().to(device)
optimizer_ft = optim.SGD(inception.parameters(), lr=0.0001, momentum=0.9)

inception_best = train_model_inception(inception, criterion, optimizer_ft,
                                       num_epochs=20, log_file_name='inception_60_0dot0001')

print('-' * 10)
evaluate_model(inception_best, log_file_name='inception_60_0dot0001')

```

Downloading:

"https://download.pytorch.org/models/inception_v3_google-1a9a5a14.pth" to
 /root/.cache/torch/checkpoints/inception_v3_google-1a9a5a14.pth

HBox(children=(FloatProgress(value=0.0, max=108857766.0), HTML(value='')))

Epoch 0/19

train Loss: 0.8416 Acc: 0.0123

val Loss: 0.5818 Acc: 0.0758

Epoch 1/19

train Loss: 0.8212 Acc: 0.0283

val Loss: 0.5584 Acc: 0.2350

Epoch 2/19

train Loss: 0.8009 Acc: 0.0683

val Loss: 0.5254 Acc: 0.4383

Epoch 3/19

train Loss: 0.7730 Acc: 0.1233

val Loss: 0.4849 Acc: 0.5383

Epoch 4/19

train Loss: 0.7474 Acc: 0.1715

val Loss: 0.4345 Acc: 0.6142

Epoch 5/19

train Loss: 0.7174 Acc: 0.2110

val Loss: 0.3832 Acc: 0.6600

Epoch 6/19

train Loss: 0.6923 Acc: 0.2290

val Loss: 0.3409 Acc: 0.6617

Epoch 7/19

train Loss: 0.6648 Acc: 0.2667

val Loss: 0.2844 Acc: 0.7042

Epoch 8/19

train Loss: 0.6409 Acc: 0.2742

val Loss: 0.2521 Acc: 0.7108

Epoch 9/19

train Loss: 0.6159 Acc: 0.3050

val Loss: 0.2178 Acc: 0.7300

Epoch 10/19

train Loss: 0.5935 Acc: 0.3238

val Loss: 0.1984 Acc: 0.7358

Epoch 11/19

train Loss: 0.5728 Acc: 0.3465

val Loss: 0.1715 Acc: 0.7517

Epoch 12/19

train Loss: 0.5541 Acc: 0.3575

val Loss: 0.1517 Acc: 0.7575

Epoch 13/19

train Loss: 0.5366 Acc: 0.3665

val Loss: 0.1401 Acc: 0.7700

Epoch 14/19

train Loss: 0.5252 Acc: 0.3744

val Loss: 0.1285 Acc: 0.7817

Epoch 15/19

train Loss: 0.5170 Acc: 0.3804

val Loss: 0.1170 Acc: 0.7917

Epoch 16/19

train Loss: 0.5012 Acc: 0.4052

val Loss: 0.1141 Acc: 0.7775

Epoch 17/19

train Loss: 0.4929 Acc: 0.4021

val Loss: 0.0989 Acc: 0.8017

Epoch 18/19

train Loss: 0.4802 Acc: 0.4142

val Loss: 0.0980 Acc: 0.7900

Epoch 19/19

train Loss: 0.4763 Acc: 0.4102

val Loss: 0.0964 Acc: 0.7933

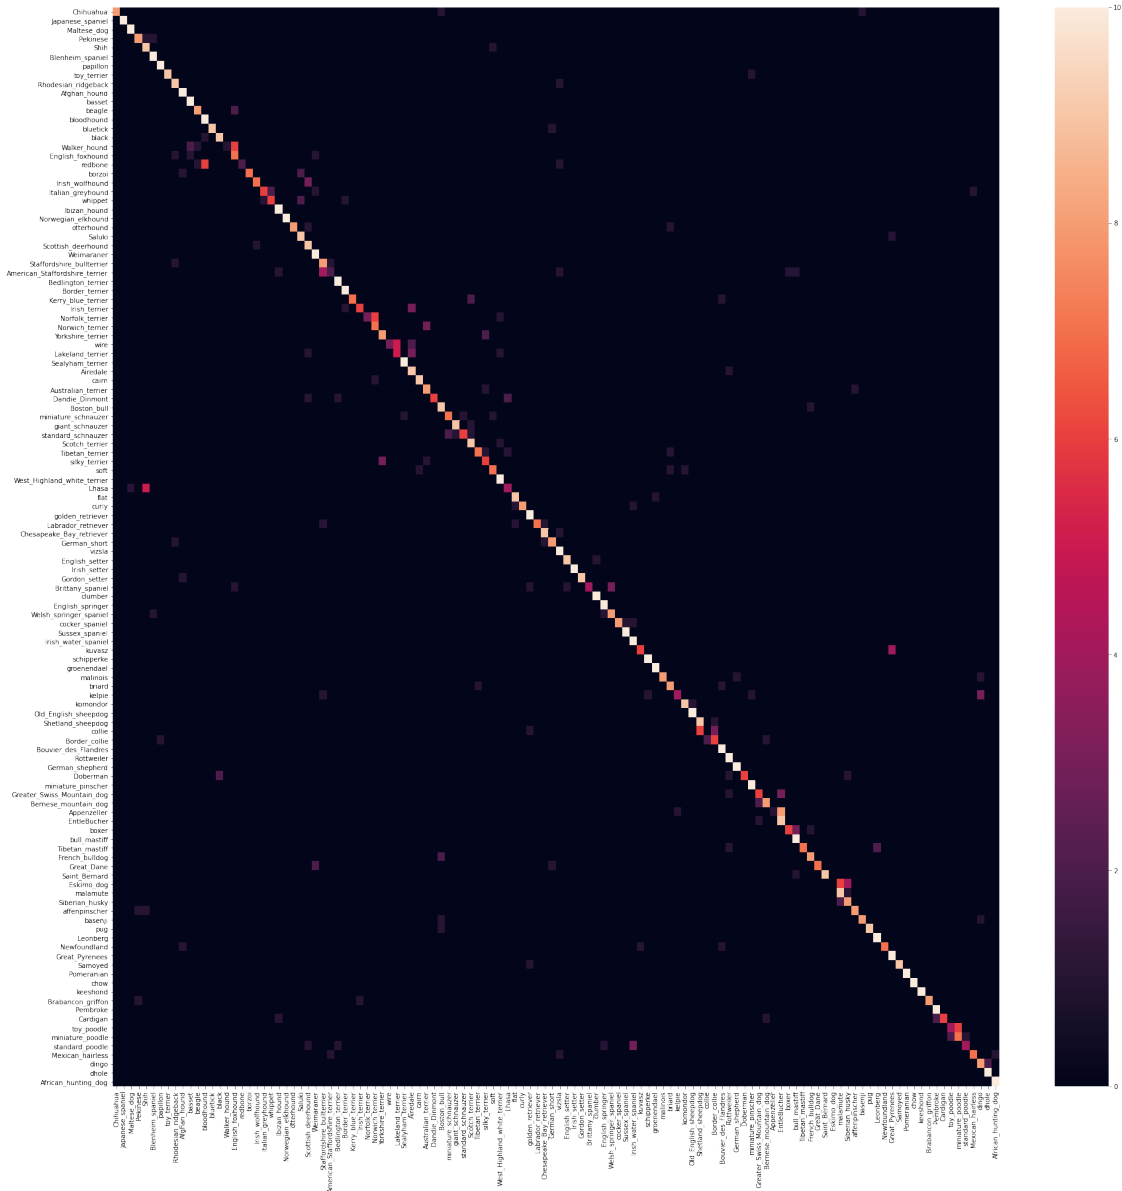
Training complete in 31m 30s

Best val Acc: 0.801667

test Loss: 0.0990 Acc: 0.7867

```
[0]: cm_inception = cm(inception_best)
df_cm = pd.DataFrame(cm_inception.numpy(), index = [i for i in my_class_names],
                    columns = [i for i in my_class_names])
plt.figure(figsize = (30,30))
sn.heatmap(df_cm)
```

```
plt.savefig('confusion_matrix_inception')
```



```
[0]: torch.save(inception_best.state_dict(), 'inception_60')
```

```
[0]: inception = models.inception_v3()

inception.AuxLogits.fc = nn.Linear(768, 120)
inception.fc = nn.Linear(2048, 120)

inception.load_state_dict(torch.load('inception_60'))
inception.eval()
```

```

[0]: Inception3(
    (Conv2d_1a_3x3): BasicConv2d(
        (conv): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), bias=False)
        (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (Conv2d_2a_3x3): BasicConv2d(
        (conv): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (Conv2d_2b_3x3): BasicConv2d(
        (conv): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (Conv2d_3b_1x1): BasicConv2d(
        (conv): Conv2d(64, 80, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(80, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (Conv2d_4a_3x3): BasicConv2d(
        (conv): Conv2d(80, 192, kernel_size=(3, 3), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (Mixed_5b): InceptionA(
        (branch1x1): BasicConv2d(
            (conv): Conv2d(192, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch5x5_1): BasicConv2d(
            (conv): Conv2d(192, 48, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch5x5_2): BasicConv2d(
            (conv): Conv2d(48, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2),
bias=False)
            (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3dbl_1): BasicConv2d(
            (conv): Conv2d(192, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
        (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_3): BasicConv2d(
        (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(192, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(32, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (Mixed_5c): InceptionA(
        (branch1x1): BasicConv2d(
            (conv): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch5x5_1): BasicConv2d(
            (conv): Conv2d(256, 48, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch5x5_2): BasicConv2d(
            (conv): Conv2d(48, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2),
bias=False)
            (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3dbl_1): BasicConv2d(
            (conv): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3dbl_2): BasicConv2d(
            (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
            (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)

```

```

    )
    (branch3x3dbl_3): BasicConv2d(
        (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(Mixed_5d): InceptionA(
    (branch1x1): BasicConv2d(
        (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch5x5_1): BasicConv2d(
        (conv): Conv2d(288, 48, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(48, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch5x5_2): BasicConv2d(
        (conv): Conv2d(48, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2),
bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_1): BasicConv2d(
        (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
        (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_3): BasicConv2d(
        (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
)

```

```

        (branch_pool): BasicConv2d(
          (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (Mixed_6a): InceptionB(
      (branch3x3): BasicConv2d(
        (conv): Conv2d(288, 384, kernel_size=(3, 3), stride=(2, 2), bias=False)
        (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch3x3dbl_1): BasicConv2d(
        (conv): Conv2d(288, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(64, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch3x3dbl_2): BasicConv2d(
        (conv): Conv2d(64, 96, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch3x3dbl_3): BasicConv2d(
        (conv): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (Mixed_6b): InceptionC(
      (branch1x1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch7x7_1): BasicConv2d(
        (conv): Conv2d(768, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch7x7_2): BasicConv2d(
        (conv): Conv2d(128, 128, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
      )
      (branch7x7_3): BasicConv2d(

```

```

        (conv): Conv2d(128, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_1): BasicConv2d(
        (conv): Conv2d(768, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_2): BasicConv2d(
        (conv): Conv2d(128, 128, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_3): BasicConv2d(
        (conv): Conv2d(128, 128, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_4): BasicConv2d(
        (conv): Conv2d(128, 128, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_5): BasicConv2d(
        (conv): Conv2d(128, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(Mixed_6c): InceptionC(
    (branch1x1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7_1): BasicConv2d(

```



```

        (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7_2): BasicConv2d(
        (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7_3): BasicConv2d(
        (conv): Conv2d(160, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_1): BasicConv2d(
        (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_2): BasicConv2d(
        (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_3): BasicConv2d(
        (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_4): BasicConv2d(
        (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_5): BasicConv2d(
        (conv): Conv2d(160, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(Mixed_6d): InceptionC(
  (branch1x1): BasicConv2d(
    (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (branch7x7_1): BasicConv2d(
    (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (branch7x7_2): BasicConv2d(
    (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (branch7x7_3): BasicConv2d(
    (conv): Conv2d(160, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (branch7x7dbl_1): BasicConv2d(
    (conv): Conv2d(768, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (branch7x7dbl_2): BasicConv2d(
    (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (branch7x7dbl_3): BasicConv2d(
    (conv): Conv2d(160, 160, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
    (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (branch7x7dbl_4): BasicConv2d(
    (conv): Conv2d(160, 160, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)

```

```

        (bn): BatchNorm2d(160, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_5): BasicConv2d(
        (conv): Conv2d(160, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(Mixed_6e): InceptionC(
    (branch1x1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7_1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7_2): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7_3): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_2): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    (branch7x7dbl_3): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_4): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7dbl_5): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (AuxLogits): InceptionAux(
        (conv0): BasicConv2d(
            (conv): Conv2d(768, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(128, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (conv1): BasicConv2d(
            (conv): Conv2d(128, 768, kernel_size=(5, 5), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(768, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (fc): Linear(in_features=768, out_features=120, bias=True)
    )
    (Mixed_7a): InceptionD(
        (branch3x3_1): BasicConv2d(
            (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3_2): BasicConv2d(
            (conv): Conv2d(192, 320, kernel_size=(3, 3), stride=(2, 2), bias=False)

```

```

        (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7x3_1): BasicConv2d(
        (conv): Conv2d(768, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7x3_2): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(1, 7), stride=(1, 1), padding=(0,
3), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7x3_3): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(7, 1), stride=(1, 1), padding=(3,
0), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch7x7x3_4): BasicConv2d(
        (conv): Conv2d(192, 192, kernel_size=(3, 3), stride=(2, 2), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (Mixed_7b): InceptionE(
        (branch1x1): BasicConv2d(
            (conv): Conv2d(1280, 320, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3_1): BasicConv2d(
            (conv): Conv2d(1280, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3_2a): BasicConv2d(
            (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0,
1), bias=False)
            (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3_2b): BasicConv2d(
            (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1,
0), bias=False)
            (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    )
    (branch3x3dbl_1): BasicConv2d(
        (conv): Conv2d(1280, 448, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(448, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
        (conv): Conv2d(448, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_3a): BasicConv2d(
        (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0,
1), bias=False)
        (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_3b): BasicConv2d(
        (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1,
0), bias=False)
        (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
        (conv): Conv2d(1280, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (Mixed_7c): InceptionE(
        (branch1x1): BasicConv2d(
            (conv): Conv2d(2048, 320, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(320, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3_1): BasicConv2d(
            (conv): Conv2d(2048, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (branch3x3_2a): BasicConv2d(
            (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0,
1), bias=False)
            (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)

```

```

    )
    (branch3x3_2b): BasicConv2d(
      (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1,
0), bias=False)
      (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_1): BasicConv2d(
      (conv): Conv2d(2048, 448, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(448, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_2): BasicConv2d(
      (conv): Conv2d(448, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_3a): BasicConv2d(
      (conv): Conv2d(384, 384, kernel_size=(1, 3), stride=(1, 1), padding=(0,
1), bias=False)
      (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch3x3dbl_3b): BasicConv2d(
      (conv): Conv2d(384, 384, kernel_size=(3, 1), stride=(1, 1), padding=(1,
0), bias=False)
      (bn): BatchNorm2d(384, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (branch_pool): BasicConv2d(
      (conv): Conv2d(2048, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn): BatchNorm2d(192, eps=0.001, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (fc): Linear(in_features=2048, out_features=120, bias=True)
  )

```

9 Finally showing some test image results

```

[0]: from PIL import Image
from torch.autograd import Variable
imsize = 299
loader = transforms.Compose([transforms.Resize(imsize), transforms.ToTensor()])

```

```

def image_loader(image_name):
    """load image, returns cuda tensor"""
    image = Image.open(image_name)
    image = loader(image).float()
    image = Variable(image, requires_grad=True)
    image = image.unsqueeze(0) #this is for VGG, may not be needed for ResNet
    return image.cuda() #assumes that you're using GPU

# image = image_loader('/content/images120_60_test/n02086240-Shih-Tzu/
↳n02086240_1725.jpg')
# image = image_loader('/content/german_shepherd.jpg')

# outputs = inception_best(image)
# _, pred = torch.max(outputs, 1)
# print(my_class_names[pred])

```

```

[0]: # show image
image_name = '/content/beagle.jpg'
plt.figure(figsize=(5,10))
plt.imshow(Image.open(image_name))
plt.axis('off')
plt.show()

# prediction
image = image_loader(image_name)
outputs = inception_best(image)
_, pred = torch.max(outputs, 1)
print('Predict as --- ', my_class_names[pred])

```




Predict as --- beagle

```
[0]: # show image
image_name = '/content/german_shepherd.jpg'
plt.figure(figsize=(5,10))
plt.imshow(Image.open(image_name))
plt.axis('off')
plt.show()

# prediction
image = image_loader(image_name)
outputs = inception_best(image)
_, pred = torch.max(outputs, 1)
print('Predict as --- ', my_class_names[pred])
```

£2000 REWARD

PLEASE HELP US FIND SAN
OUR GERMAN SHEPHERD DOG



****STILL MISSING****

San our German Shepherd Male.

He has Ear Tattoo Identification.

Was last seen at Home

Sowerby Bridge/Ripponden/Cragg Vale
area of West Yorkshire on Saturday 23rd
September 2017 at 5pm. Feared Stolen.

Please contact his owners on
07727 424624.

Predict as --- German_shepherd

```
[0]: # show image
image_name = '/content/gsm�.jpg'
plt.figure(figsize=(5,10))
plt.imshow(Image.open(image_name))
plt.axis('off')
plt.show()

# prediction
```

```
image = image_loader(image_name)
outputs = inception_best(image)
_, pred = torch.max(outputs, 1)
print('Predict as --- ', my_class_names[pred])
```



Predict as --- Greater_Swiss_Mountain_dog

```
[0]: class_names
```

```
[0]: ['n02085620-Chihuahua',
      'n02085782-Japanese_spaniel',
      'n02085936-Maltese_dog',
      'n02086079-Pekinese',
      'n02086240-Shih-Tzu',
      'n02086646-Blenheim_spaniel',
      'n02086910-papillon',
      'n02087046-toy_terrier',
      'n02087394-Rhodesian_ridgeback',
      'n02088094-Afghan_hound',
      'n02088238-basset',
      'n02088364-beagle',
      'n02088466-bloodhound',
      'n02088632-bluetick',
      'n02089078-black-and-tan_coonhound',
      'n02089867-Walker_hound',
      'n02089973-English_foxhound',
      'n02090379-redbone',
      'n02090622-borzoi',
      'n02090721-Irish_wolfhound',
```

'n02091032-Italian_greyhound',
'n02091134-whippet',
'n02091244-Ibizan_hound',
'n02091467-Norwegian_elkhound',
'n02091635-otterhound',
'n02091831-Saluki',
'n02092002-Scottish_deerhound',
'n02092339-Weimaraner',
'n02093256-Staffordshire_bullterrier',
'n02093428-American_Staffordshire_terrier',
'n02093647-Bedlington_terrier',
'n02093754-Border_terrier',
'n02093859-Kerry_blue_terrier',
'n02093991-Irish_terrier',
'n02094114-Norfolk_terrier',
'n02094258-Norwich_terrier',
'n02094433-Yorkshire_terrier',
'n02095314-wire-haired_fox_terrier',
'n02095570-Lakeland_terrier',
'n02095889-Sealyham_terrier',
'n02096051-Airedale',
'n02096177-cairn',
'n02096294-Australian_terrier',
'n02096437-Dandie_Dinmont',
'n02096585-Boston_bull',
'n02097047-miniature_schnauzer',
'n02097130-giant_schnauzer',
'n02097209-standard_schnauzer',
'n02097298-Scotch_terrier',
'n02097474-Tibetan_terrier',
'n02097658-silky_terrier',
'n02098105-soft-coated_wheaten_terrier',
'n02098286-West_Highland_white_terrier',
'n02098413-Lhasa',
'n02099267-flat-coated_retriever',
'n02099429-curly-coated_retriever',
'n02099601-golden_retriever',
'n02099712-Labrador_retriever',
'n02099849-Chesapeake_Bay_retriever',
'n02100236-German_short-haired_pointer',
'n02100583-vizsla',
'n02100735-English_setter',
'n02100877-Irish_setter',
'n02101006-Gordon_setter',
'n02101388-Brittany_spaniel',
'n02101556-clumber',
'n02102040-English_springer',

'n02102177-Welsh_springer_spaniel',
'n02102318-cocker_spaniel',
'n02102480-Sussex_spaniel',
'n02102973-Irish_water_spaniel',
'n02104029-kuvasz',
'n02104365-schipperke',
'n02105056-groenendael',
'n02105162-malinois',
'n02105251-briard',
'n02105412-kelpie',
'n02105505-komondor',
'n02105641-Old_English_sheepdog',
'n02105855-Shetland_sheepdog',
'n02106030-collie',
'n02106166-Border_collie',
'n02106382-Bouvier_des_Flandres',
'n02106550-Rottweiler',
'n02106662-German_shepherd',
'n02107142-Doberman',
'n02107312-miniature_pinscher',
'n02107574-Greater_Swiss_Mountain_dog',
'n02107683-Bernese_mountain_dog',
'n02107908-Appenzeller',
'n02108000-EntleBucher',
'n02108089-boxer',
'n02108422-bull_mastiff',
'n02108551-Tibetan_mastiff',
'n02108915-French_bulldog',
'n02109047-Great_Dane',
'n02109525-Saint_Bernard',
'n02109961-Eskimo_dog',
'n02110063-malamute',
'n02110185-Siberian_husky',
'n02110627-affenpinscher',
'n02110806-basenji',
'n02110958-pug',
'n02111129-Leonberg',
'n02111277-Newfoundland',
'n02111500-Great_Pyrenees',
'n02111889-Samoyed',
'n02112018-Pomeranian',
'n02112137-chow',
'n02112350-keeshond',
'n02112706-Brabancon_griffon',
'n02113023-Pembroke',
'n02113186-Cardigan',
'n02113624-toy_poodle',

```
'n02113712-miniature_poodle',  
'n02113799-standard_poodle',  
'n02113978-Mexican_hairless',  
'n02115641-dingo',  
'n02115913-dhole',  
'n02116738-African_hunting_dog']
```

```
[0]: with open('log/inception_60_0dot0001', 'r') as inFile:  
      for line in inFile.readlines():  
          print(line.split(' '))
```

```
['0', '0.8416', '0.0123', 'train\n']  
['0', '0.5818', '0.0758', 'val\n']  
['1', '0.8212', '0.0283', 'train\n']  
['1', '0.5584', '0.2350', 'val\n']  
['2', '0.8009', '0.0683', 'train\n']  
['2', '0.5254', '0.4383', 'val\n']  
['3', '0.7730', '0.1233', 'train\n']  
['3', '0.4849', '0.5383', 'val\n']  
['4', '0.7474', '0.1715', 'train\n']  
['4', '0.4345', '0.6142', 'val\n']  
['5', '0.7174', '0.2110', 'train\n']  
['5', '0.3832', '0.6600', 'val\n']  
['6', '0.6923', '0.2290', 'train\n']  
['6', '0.3409', '0.6617', 'val\n']  
['7', '0.6648', '0.2667', 'train\n']  
['7', '0.2844', '0.7042', 'val\n']  
['8', '0.6409', '0.2742', 'train\n']  
['8', '0.2521', '0.7108', 'val\n']  
['9', '0.6159', '0.3050', 'train\n']  
['9', '0.2178', '0.7300', 'val\n']  
['10', '0.5935', '0.3238', 'train\n']  
['10', '0.1984', '0.7358', 'val\n']  
['11', '0.5728', '0.3465', 'train\n']  
['11', '0.1715', '0.7517', 'val\n']  
['12', '0.5541', '0.3575', 'train\n']  
['12', '0.1517', '0.7575', 'val\n']  
['13', '0.5366', '0.3665', 'train\n']  
['13', '0.1401', '0.7700', 'val\n']  
['14', '0.5252', '0.3744', 'train\n']  
['14', '0.1285', '0.7817', 'val\n']  
['15', '0.5170', '0.3804', 'train\n']  
['15', '0.1170', '0.7917', 'val\n']  
['16', '0.5012', '0.4052', 'train\n']  
['16', '0.1141', '0.7775', 'val\n']  
['17', '0.4929', '0.4021', 'train\n']  
['17', '0.0989', '0.8017', 'val\n']
```

```
['18', '0.4802', '0.4142', 'train\n']  
['18', '0.0980', '0.7900', 'val\n']  
['19', '0.4763', '0.4102', 'train\n']  
['19', '0.0964', '0.7933', 'val\n']  
['test', '0.0990', '0.7867', 'test\n']
```

```
[0]: !zip -r log2.zip log
```

```
adding: log/ (stored 0%)  
adding: log/resnet_40_0dot0001 (deflated 61%)  
adding: log/resnet_20_0dot001 (deflated 61%)  
adding: log/resnet_20_0dot0001 (deflated 61%)  
adding: log/.ipynb_checkpoints/ (stored 0%)  
adding: log/resnet_40_0dot001 (deflated 61%)  
adding: log/vgg_40_0dot0001 (deflated 62%)  
adding: log/inception_40_0dot0001 (deflated 60%)  
adding: log/vgg_20_0dot0001 (deflated 61%)  
adding: log/inception_20_0dot0001 (deflated 60%)
```

```
[0]:
```